

Chloe Tu
ITAI 1378
September 23, 2025

L05 Reflective Journal on the Chihuahua vs Muffin CNN Lab

Introduction

My experience with the CNN-based classification task was both challenging and incredibly rewarding. Building on the foundation of the traditional neural network from the previous workshop, this lab gave me a deeper understanding of how powerful deep learning can be. The CNN architecture presented a new way to approach image classification. It showed me how a model can be designed to truly "see" and interpret visual data. This journal will reflect on my journey through the lab, from understanding the architecture to facing the challenges and considering the broader implications of this technology.

1. CNN Architecture

My experience with the CNN architecture was a significant step up from the last lab. CNN uses special layers like the convolutional layer and the pooling layer. These layers are different from the traditional neural network I used before. The convolutional layer works like a filter, moving over the image to find features such as edges or shapes. The pooling layer then reduces the size of the data which helps to make the model more efficient. This process allows the network to learn important visual characteristics in a structured way. In contrast, the traditional neural network from the previous workshop took the entire image and flattened it into one long line of pixels. This meant it lost all the spatial relationships between the pixels. CNN's design is much better for processing images because it keeps this spatial information. As we learned in L05 CNN and L04 Traditional Neural Network, this fundamental difference is what makes a CNN so effective for image classification tasks. It is able to learn complex features that a traditional network simply cannot see.

2. Model Performance

The model's performance was very good. The accuracy was high on the test data. I was able to observe how well it is classified as chihuahuas and muffins. Most of the time it got the classification right. This showed me that CNN was a very effective tool for this specific problem. However, I did notice some interesting misclassification patterns. The model sometimes confused a chihuahua with a muffin. This probably happened when the chihuahua was a light brown color and looked like the muffin's shape and texture. Or maybe the lighting in the picture made it hard for the model to tell the difference. This shows that the model still has trouble with certain visual similarities, even with its high overall accuracy. It highlights the importance of having diverse and varied data to train the model well.

3. Comparison

When I compare this CNN model to the traditional neural network from the last lab, the difference in performance is clear. CNN performed much better. Its ability to find and learn from image features gave it a big advantage. The traditional neural network could not do this as well. It resulted in lower accuracy. The traditional model's flattened input made it less effective for tasks that require an understanding of image structure. In terms of training time, CNN took longer. It has more layers and a more complex structure, so the computer needed more time to train it. The extra time was worth it, though. CNN's performance was far superior for this image task. The increased computational cost is a good trade-off for the higher accuracy and better performance on image-based tasks. The time investment paid off with a more robust and accurate model.

4. Challenges and Solutions

The main challenge I faced was making sure the model did not just memorize the training data. This problem is called overfitting. At first, my model might have been doing too well on the training set but not so well on new pictures. This can be a common issue when building a deep learning model. It is something I had to actively work to correct during the lab. To fix this, I used methods to help the model generalize. I added a dropout layer to the model. This layer randomly turns off some of the neurons during training, which forces the network to find other ways to make predictions. I also used a technique to change the training pictures slightly which is called data augmentation. This made the model see new versions of the images. By doing these things, I was able to make the model smarter and better at classifying pictures it had never seen before.

5. Real-World Applications

This type of image classification has many exciting uses in the real world. For example, it could be used in medical imaging to help doctors find diseases or tumors more quickly. It could also be used in self-driving cars to help them identify pedestrians, signs, and other vehicles on the road. This allows the car to make informed decisions and navigate safely. Another application is in retail, where it could automatically sort products and manage inventory by identifying different items. Companies could also use it to moderate content online, flagging inappropriate images. Technology is very useful and has the potential to solve many different problems. This is just a small sample of the countless ways this technology can be applied.

6. Ethical Considerations

There are important ethical considerations with these kinds of models. One concern is bias in the data. If the model is trained on a biased dataset, it might not work well for everyone. For example, a facial recognition model trained on a limited set of ethnicities might perform poorly on others. This could lead to unfair or inaccurate results for certain groups of people. Another big concern is privacy. Using these models for things like facial recognition raises questions about privacy and surveillance. Developers need to be careful about how these models are used and how the data is collected. It is important to make sure the technology is used for good things and does not harm people. I believe we must always think about these issues as we create and use these powerful models. We have a responsibility to use this technology in a way that benefits all of society.

Conclusion

Overall, this lab was an eye-opening experience that went beyond simple coding. It taught me about the practical and ethical aspects of deep learning. I learned that designing a good model is not just about writing code. It is about understanding the data, fixing problems, and thinking about the real-world impact. CNN proved to be a powerful tool for image classification, far more effective than the traditional neural network. I am excited to continue exploring this field and using these skills to create meaningful solutions.

References

- Gill, Navdeep Singh. “Dark Web Monitoring with SOC Automation.” Xenonstack.com, 28 Nov. 2024, <https://www.xenonstack.com/blog/ethical-considerations-in-computer-vision>.
- AWS. “What Is Overfitting? - Overfitting - AWS.” Amazon Web Services, Inc., 2024, <https://aws.amazon.com/what-is/overfitting/>.
- “Convolutional Neural Network (CNN).” NVIDIA Developer, 23 Apr. 2018, <https://developer.nvidia.com/discover/convolutional-neural-network>.