Chloe Tu
ITAI 2373
September 28, 2025

Reflection Report: Mapping the Neural Network Brain

1. Introduction

This report reflects on the process of creating a knowledge graph about neural networks as part of the "A06 Mapping the Neural Network Brain" project. Before starting, my understanding of AI was a scattered list of terms and abstract ideas that I struggled to connect into a coherent whole. The goal of this project was to transform that fragmented knowledge by visualizing the core components and processes that define this field. This document discusses how the mapping exercise made these concepts tangible and how it revealed surprising relationships I had never considered. It also explores the utility of biological analogies in learning this material. Ultimately, this report details how my personal understanding evolved from a simple list of definitions into a cohesive, system-level perspective, capturing the key insights gained through this valuable learning experience.

2. How creating the graph helped clarify neural network concepts

The primary benefit of the mapping exercise was its power to transform abstract definitions into a structured, visual system that I could finally navigate. Creating the knowledge graph forced me to organize disparate ideas into a logical hierarchy, making their relationships clear for the first time. For instance, mapping the flow from a Single Neuron to a Perceptron and then to a Multi-Layer Perceptron / Deep Network made the progression of complexity tangible. For the first time, I could visually trace the evolution from a simple processing unit to a complex network, making the leap in capability feel earned and understandable rather than magical. Similarly, visualizing the learning process as a four-step cycle of Training Data, a Forward Pass, Backpropagation, and a Feedback Loop demystified the iterative nature of AI training. It became a clear, repeatable process instead of an opaque black box. Connecting architectural components like Hidden Layers to their purpose of extracting complex features and Activation Functions to their role in introducing non-linearity solidified their individual contributions. This process turned a collection of isolated facts into an interconnected machine where every part has a clear function. Beyond clarifying individual ideas, this process of forced organization was crucial for revealing unexpected relationships between them.

3. Surprising connections discovered between concepts

The knowledge graph format forces the creation of links between topics that are not obvious when they are learned in isolation, leading to several powerful "aha moments." The first surprising insight was the link between a biological idea and a data processing technique, specifically the description of Data augmentation as "expanded dendrites." This analogy reframed a purely technical step for increasing training data into an intuitive extension of a neuron's input mechanism, making its purpose far more memorable. Another profound connection was the parallel drawn between the electrical impulse speed in a biological Axon and the latency of inference in an artificial system. This connection was startling because it linked a microscopic biological constraint to a macro-level engineering challenge that determines the viability of technologies like self-driving cars. It reframed inference speed from a simple performance metric into a direct artificial echo of a biological imperative for rapid response. Lastly, mapping out different learning theories highlighted a key distinction: contrasting the local, unsupervised nature of Hebbian Learning with the global, supervised method of Gradient-based Learning clarified exactly how AI learning is inspired by the brain but is not a direct copy of its processes. While these surprising connections were intellectually stimulating, evaluating their overall utility was a critical next step in my learning.

4. Which biological-artificial parallels were most/least helpful
Biological analogies served a dual role throughout this learning process, acting as both a powerful teaching tool and a potential source of oversimplification. The most helpful parallels were the core structural analogies that provide an intuitive foundation for understanding a neuron's basic function. Mapping Dendrites to Inputs, Synapses to Weights, and the Axon to the Output created a memorable framework for the artificial neuron. Furthermore, the biological firing threshold was an incredibly helpful parallel for understanding the critical role of Activation Functions in introducing non-linearity. However, the analogy becomes far less useful when examining the specific mechanisms of network learning. The simple Hebbian principle of "Neurons that fire together, wire together" is conceptually distant from the complex, mathematically driven process of modern AI. The analogy fails because Hebbian Learning is a fundamentally local and unsupervised process. In stark contrast, Backpropagation is a global and supervised process that requires a system-wide error signal and knowledge of the entire network's structure to calculate gradients. This distinction clarified that AI learning isn't just inspired by the brain; it's a completely different engineering paradigm. Grappling with both the helpful and unhelpful aspects of these analogies was a key part of the learning journey.
5. My Evolving Understanding of Neural Networks evolved through the mapping process
Before this project, my knowledge of neural networks was a jumble of disconnected facts and vocabulary. The mapping process was instrumental in transforming that list into a cohesive, system-level perspective that I can now articulate. My understanding shifted from seeing concepts like Optimizers and Loss Functions as isolated terms to grasping their dynamic interplay within the Backpropagation & Weight Updates step of the learning cycle. They are not static definitions, but active components working together to minimize error. The visual link on the graph was a breakthrough in connecting problems to their solutions. I finally understood that the Perceptron's failure to solve the XOR problem wasn't just a historical footnote; it was the fundamental engineering challenge that necessitated the invention of Hidden Layers and non-linear Activation Functions in MLPs. The solution was a direct response to a critical limitation. The final stage of my understanding came when I could trace a concept from its lowest level to its highest-level application. I can now see how a fundamental component like a Weight is not just a number, but a parameter that becomes part of user/item embeddings in recsys or a key element in a diffusion model. This newfound systemic understanding provides a clear foundation for simplifying these complex ideas for a non-technical audience.
6. Suggestions for explaining neural networks to someone with no technical background
Explaining a complex technical topic simply is a significant challenge, but this project provided a clear roadmap for doing so. To explain neural networks to a novice, I would use a carefully scaffolded narrative built on the most effective analogies from the knowledge graph. The best starting point is the core Biological Neuron to Artificial Neuron parallel, describing a neuron as a simple listener with "ears" (Dendrites) and a "mouth" (Axon), where internal "volume dials" (Weights) control which sounds are most important. From there, one can introduce the learning process using the intuitive concept of "practice makes perfect" to explain the Feedback Loop of Improvement, where the network makes a guess (Predict), checks its answer (Compare), and adjusts its dials to do better next time (Update). To explain how these simple units create a powerful system, I would use the hierarchical learning example of recognizing edges → textures → objects, framing it as learning simple building blocks first to understand complex things. Finally, I would ground the entire explanation in concrete applications, such as how Face Filters learn to find landmarks or how Recommendation Systems learn a user's preferences. This step-by-step process, moving from a simple analogy to a tangible, real-world result, can make the core concepts of neural networks accessible to anyone.